

# An Overview of the Trimaran Compiler Infrastructure



# Infrastructure Goals

---

- To provide a vehicle for implementation and experimentation for state of the art research in compiler techniques for instruction-level parallel architectures.
  - Currently, the infrastructure is oriented towards Explicitly Parallel Instruction Computing (EPIC) architectures.
    - But can also support compiler research for Superscalar architectures.
  - Primarily, “back-end” compiler research
    - instruction scheduling, register allocation, and machine dependent optimizations.



# Terms and Definitions

---

- ILP (Instruction-Level Parallelism)
  - more than one operation issued per clock cycle within a single CPU
- EPIC (Explicitly Parallel Instruction Computing)
  - ILP under compiler control
    - A single instruction may contain many operations
    - Compiler determines operation dependences and specifies which operations may execute concurrently



# Terms and Definitions (cont)

---

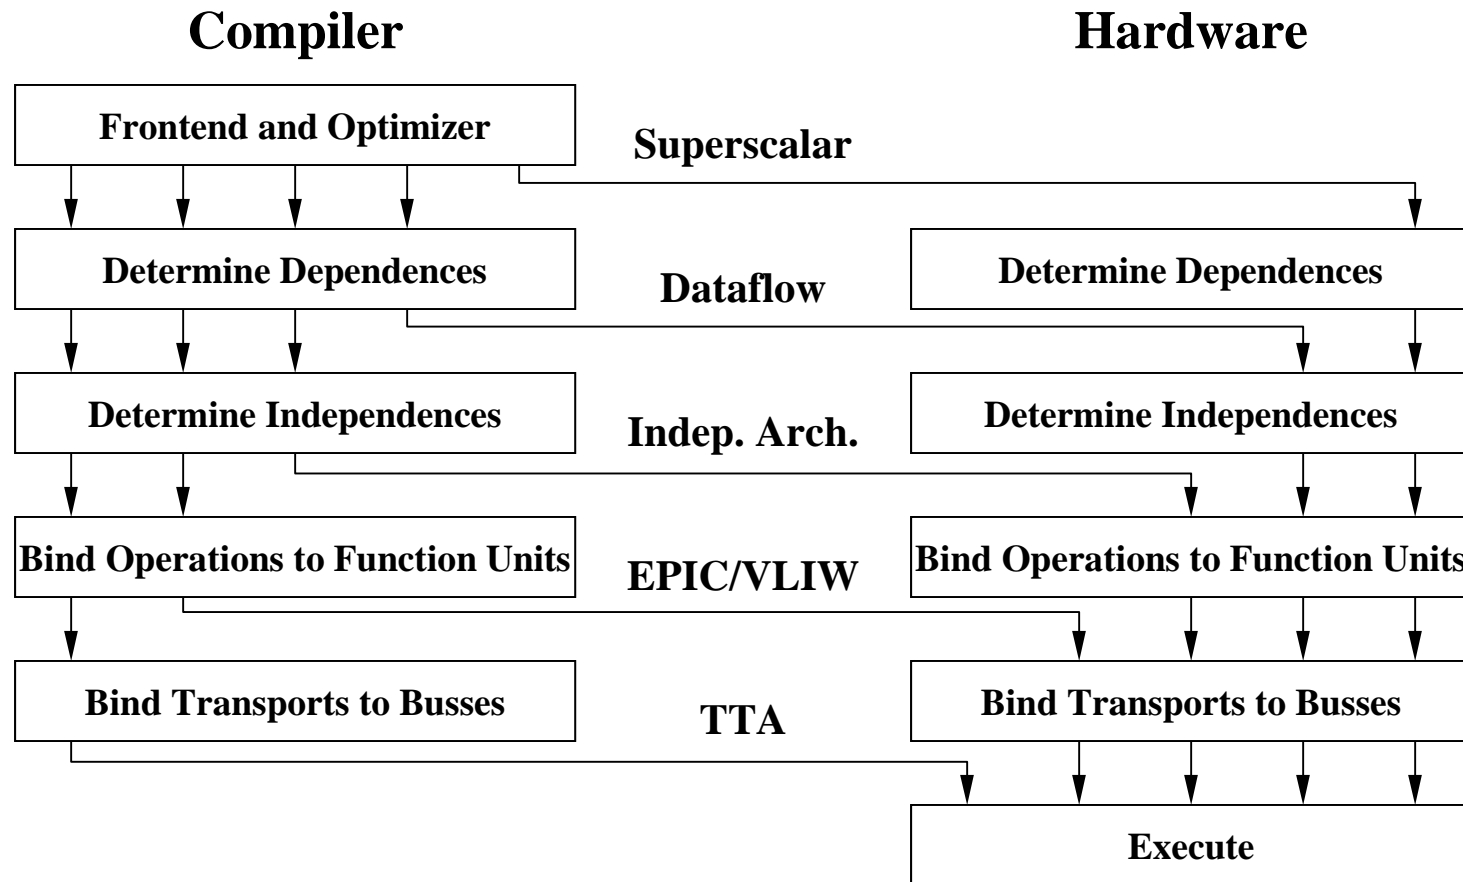
- Superscalar

- ILP under hardware control

- CPU looks ahead through the instruction stream to find instructions whose execution can be initiated in the current cycle
      - hardware finds dependences between instructions.
      - dynamic reordering of instructions



# Architectural Overview



- From “Code Generation for Transport-Triggered Architectures”, Henk Corporaal, Jan Hoogerbrugge (Delft University of Technology)



# Infrastructure Support

---

The infrastructure is comprised of the following components:

- A machine description language, HMDES, for describing ILP architectures.
- A parameterized ILP Architecture called HPL-PD
  - Current instantiation in the infrastructure is as an EPIC architecture
- A compiler front-end for C, performing parsing, type checking, and a large suite of high-level (i.e. machine independent) optimizations.
  - This is the IMPACT module (IMPACT group, University of Illinois)



# Infrastructure Support (cont)

---

- A compiler back-end, parameterized by a machine description, performing instruction scheduling, register allocation, and machine-dependent optimizations.
  - This is the Elcor module (CAR group, HP Labs)
  - Each stage of the back-end may easily be replaced or modified by a compiler researcher.



# Infrastructure Support (cont)

---

- An extensible IR (intermediate program representation)
  - Has both an internal and textual representation, with conversion routines between the two. The textual language is called Rebel.
  - Supports modern compiler techniques by representing control flow, data and control dependence, and many other attributes.
  - Easy to use in its internal representation (clear C++ object hierarchy) and textual representation (human-readable)





# Infrastructure Support (cont)

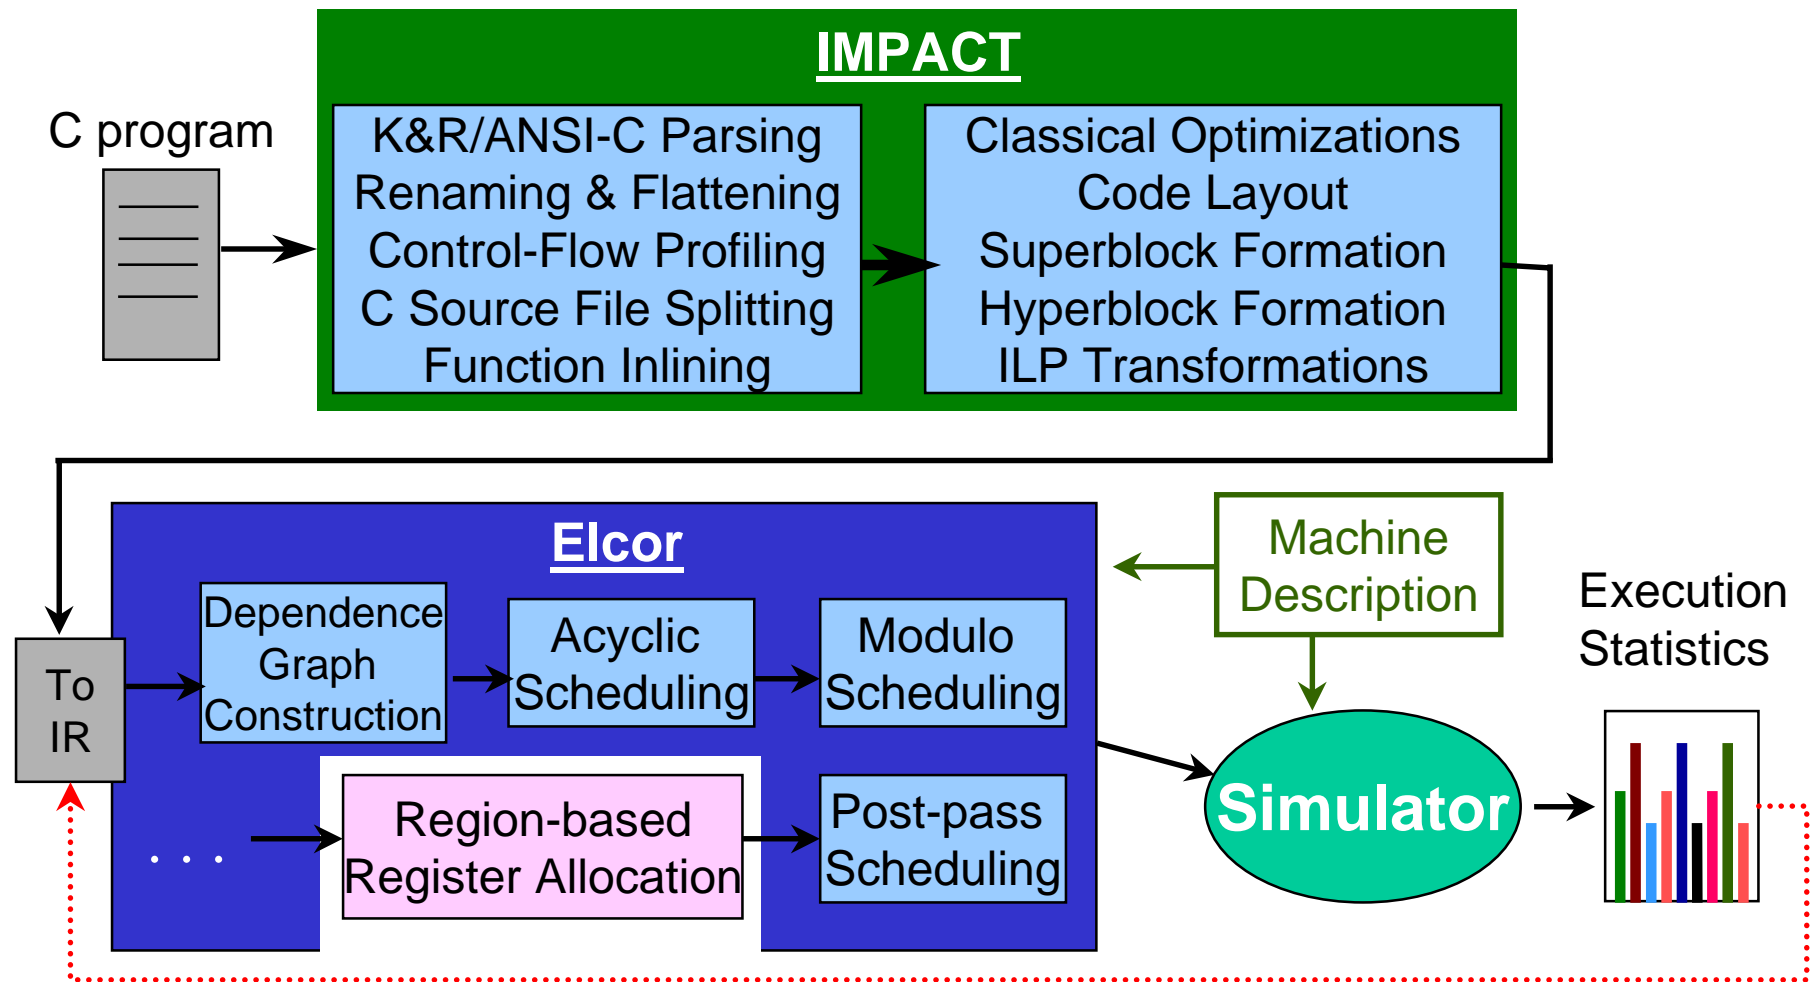
---

- A cycle-level simulator of the HPL-PD architecture which is configurable by a machine description and provides run-time information on execution time, branch frequencies, and resource utilization.
  - This information can be used for profile-driven optimizations, as well as to provide validation of new optimizations.
  - The HPL-PD simulator was implemented by the ReaCT\_ILP group at NYU.



# System Organization

A compiler researcher's view of the infrastructure:

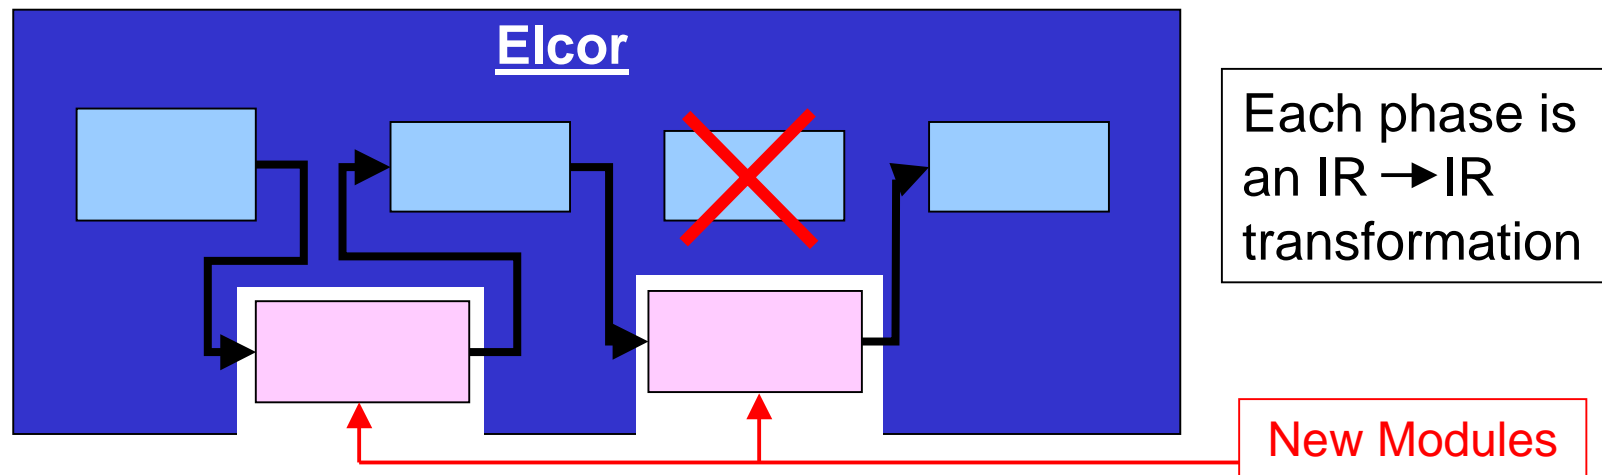




# The research process

The infrastructure is used for designing, implementing, and testing new compilation modules to be incorporated into the back end.

- These phases may augment or replace existing Elcor modules.



- New modules may be the result of research in scheduling, register allocation, program analysis, profile-driven compilation, etc.
  - For example, NYU has added a region-based register allocator.



# Why use this compiler infrastructure?

---

- It is especially geared for ILP research
- It provides a rich compilation framework
  - Parameterized ILP architecture (HPL-PD)
  - Machine description language
  - Single intermediate program representation
    - provides mechanism for representing wide range of program information
  - Cycle-level execution simulation
    - provides run-time information for profile-driven compilation



## More reasons...

---

- The framework is populated with a large number of existing compilation modules
  - provides leverage for new compiler research
  - supports meaningful experimentation, rather than simply running toy programs.
  - Full compilation and execution path already exists
- There's a commitment on our part to releasing a robust, tested, and documented software system.