

Case Study: A Region-Based Register Allocator



Overview

- The Elcor Module needed a register allocator that exploited the region-based program representation.
- A region-based register allocator was implemented.
 - Technical description to follow
- Register allocator is complete and will be part of the initial release of the system.
 - Design & implementation contributed by Kim, Gopinath, Kathail, Esfahany, and Palem.

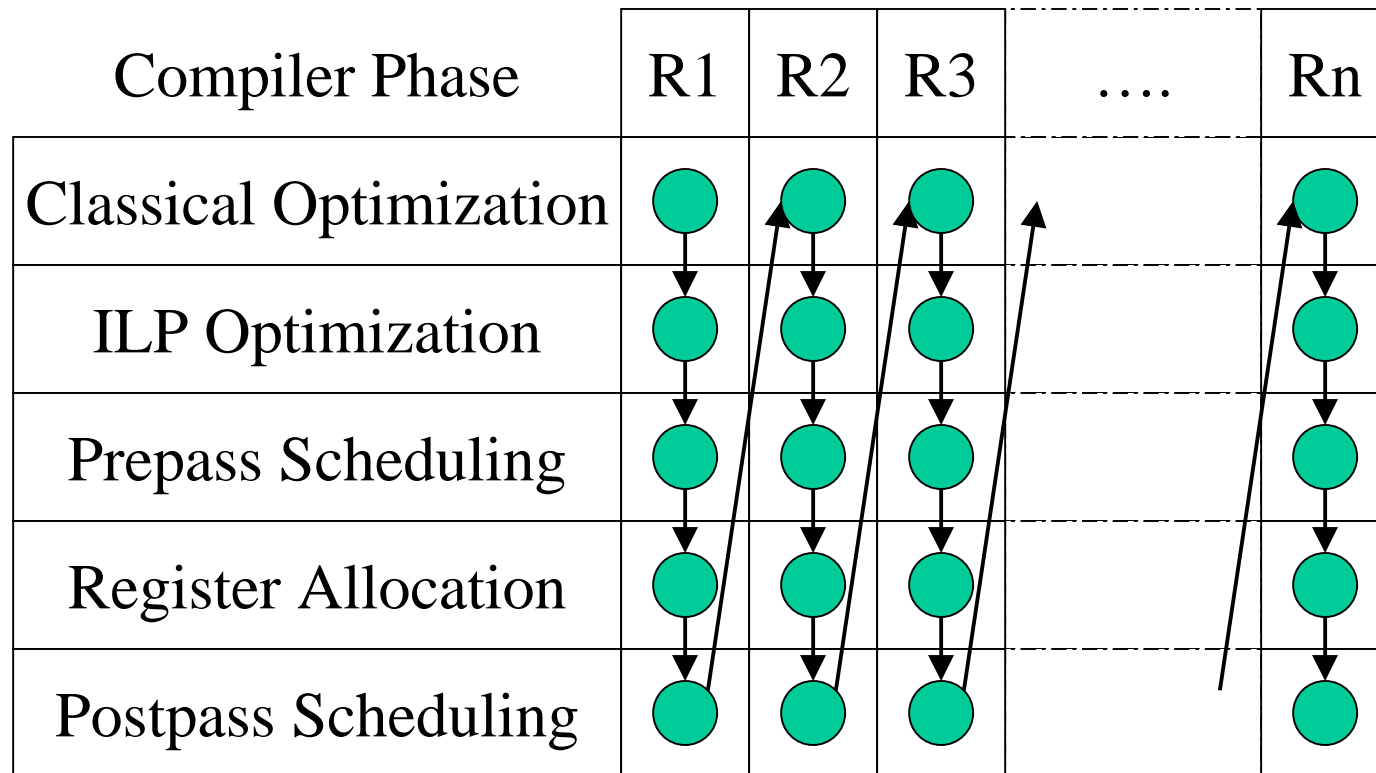


Features of the Register Allocator

- Region Based
- Priority Based Coloring
 - Chow & Hennessy's approach
 - Frequency based priority function
- Finer Grained Live Ranges
- Region Reconciliation
- Handles Predicated Instructions



Phase Order



In our actual implementation, this is performed at the granularity of procedures.



Advantages

- Compiler is in complete control over the size and contents of the compilation unit
- The size of the compilation unit is typically smaller than functions, reducing the importance of the algorithmic complexity of the applied ILP transformations.
- The use of profile information to select regions allows the compiler to select compilation units that more accurately reflect the dynamic behavior of the program.
- Reference
 - “Region Based Compilation” Richard Hanks, Wen-mei Hwu and Bob Rau




Finer Grained Live Ranges

- We use operations as the units of live range construction
 - Chow & Hennessy used basic blocks
- Problems with Basic Block-based LR construction:
 - Imprecise LR construction
 - Problem is exacerbated by Superblocks and Hyperblocks
 - Code Predication is not visible in Basic Block



Basic Block Live Range Construction

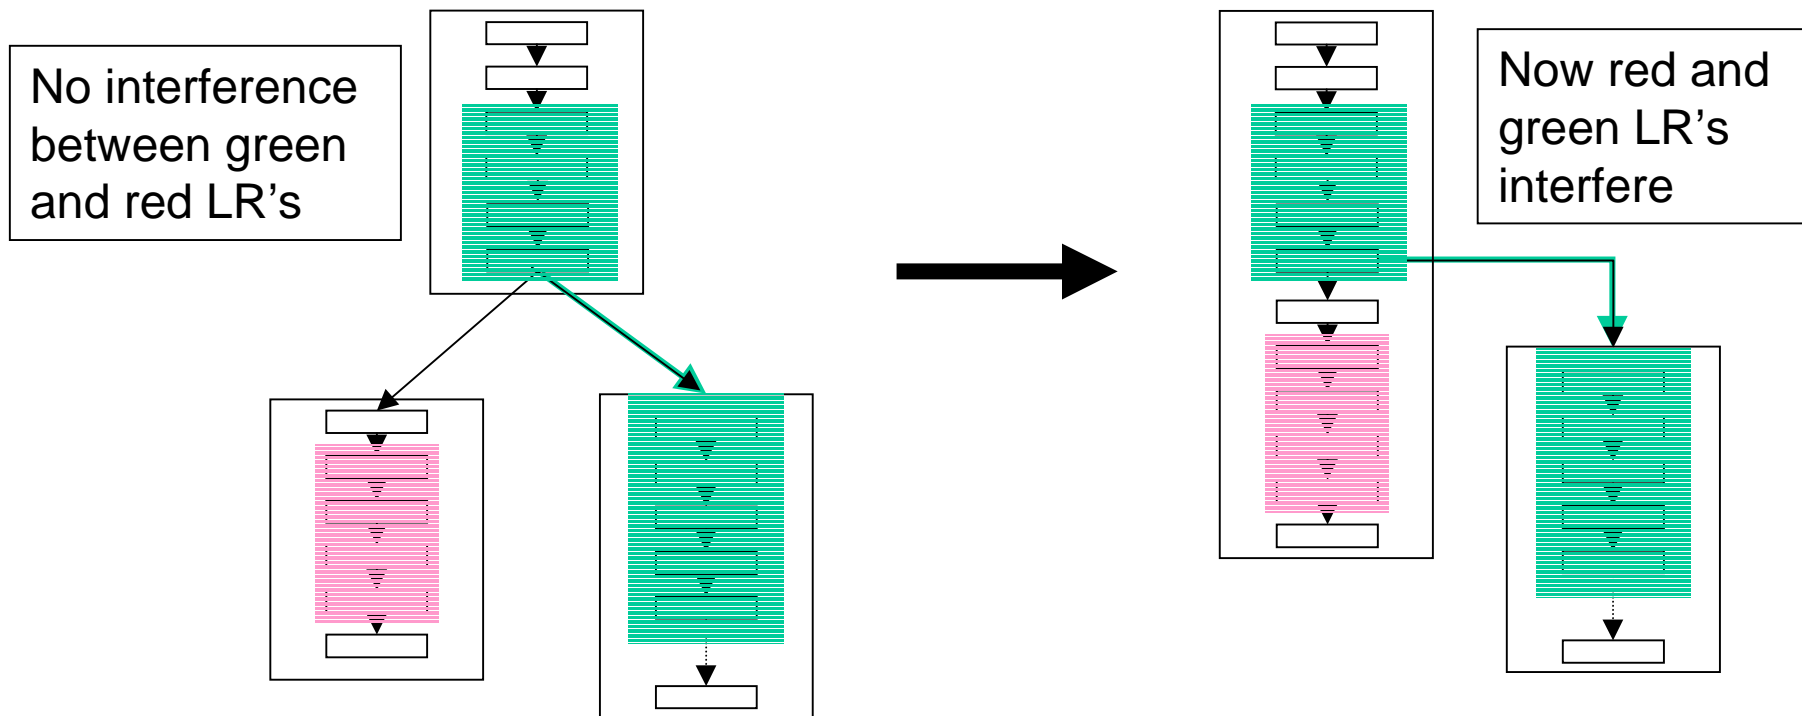
- This basic block contains two non-interfering live ranges.
- They would be marked as interfering in an interference graph using basic block live range construction.
- They would be allocated to separate physical registers.





Basic Block LR Construction (cont)

- When basic blocks are coalesced to form superblocks, the problem becomes even worse.
- New false interferences are created.





Priority Based Coloring

- Live Ranges are colored (i.e. allocated to registers) in order of their priority.
- Their priority is computed based on:
 - the frequency of execution of their containing region
 - the cost/benefit of keeping them in a register.
- Register assignment incorporates a calling cost computation
 - i.e. assignment to caller-saved vs. callee-saved registers.
 - Lueh & Gross 1997



Priority Based Coloring (cont)

- Within a single block, the priority $P(l)$ of a live range l for a virtual register v is computed by

$$P = S_l / N_l$$

where

$$S_l = \text{loadsave} * u_l + \text{storesave} * d_l$$

$$N_l = \text{Size of live range (number of operations)}$$

loadsave = cost of a load

storesave = cost of a store

u_l = number of uses of v

d_l = number of definitions of v



Priority Based Coloring (cont)

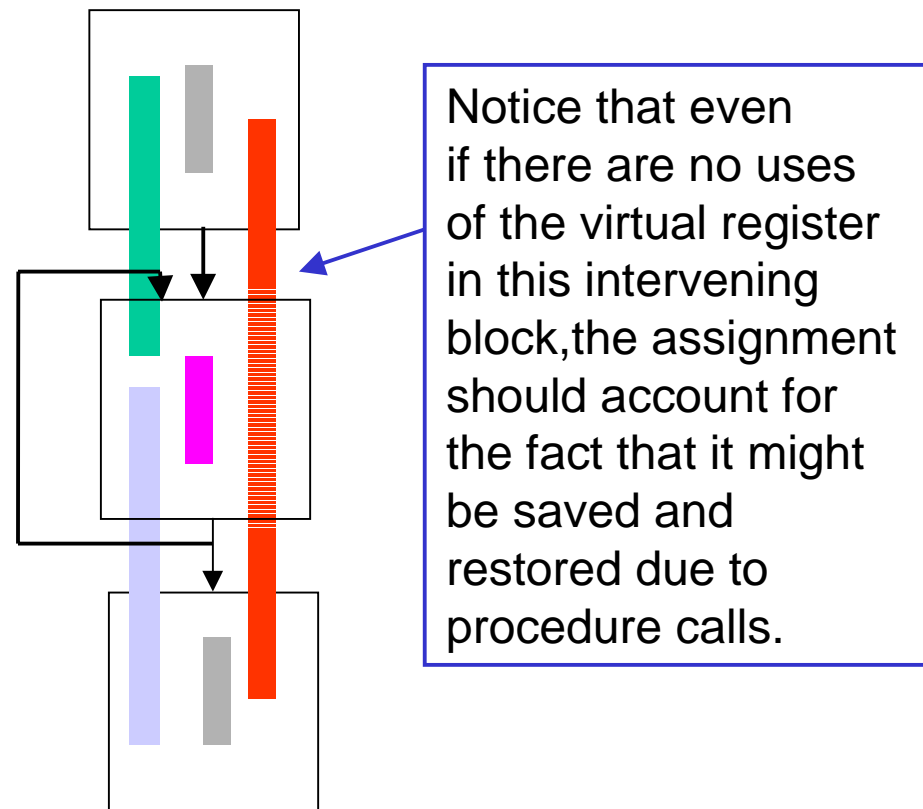
- For multi-block regions, the priority of a live range l , $P(l)$, is based on the execution frequency $F(b)$ of each block b in l and $P_b(l)$.

$$P(l) = \sum (P_b(l) * F(b))$$



Register Assignment

- Register assignment is done at the block level
 - with consideration of the use of caller-saved vs. callee-saved registers



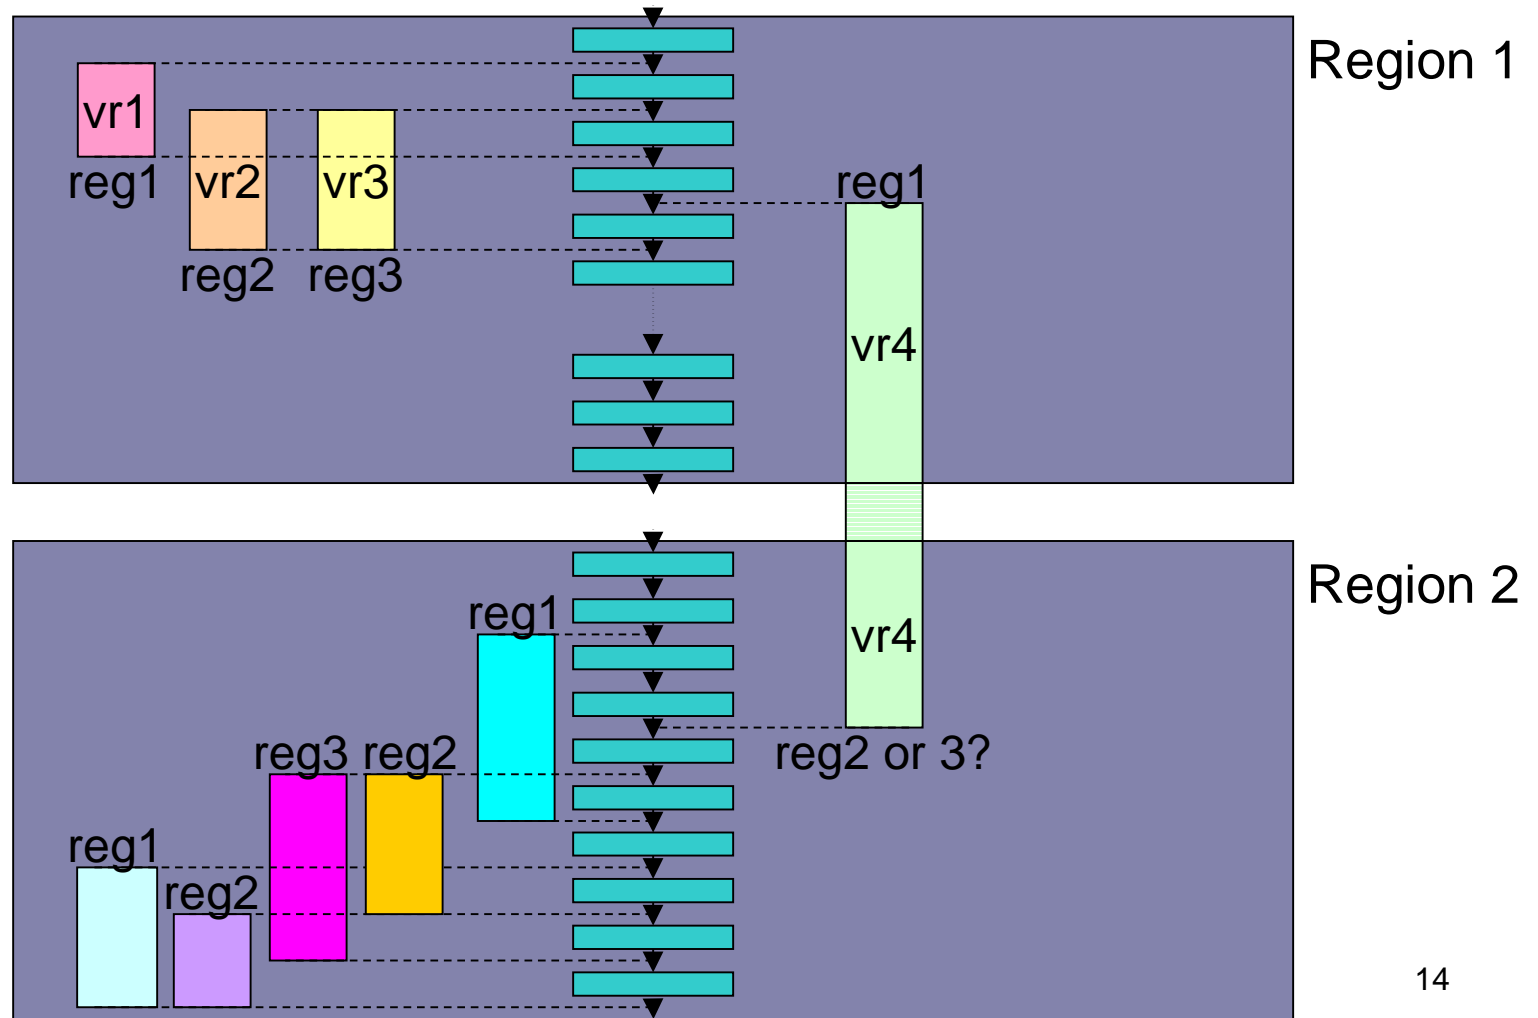


Reconciling Live Ranges

- In each region, the LR's are bound to physical registers based on their priorities.
 - The same LR in two different regions may be bound to different physical registers.
- Such an LR must be reconciled.
 - Need Patch Up code between each Region
 - may have to insert copy operations
- Reconciliation includes assigning to physical registers the delayed-bound LR's.
 - i.e. live ranges with no defs or uses in the block.



Reconciling Live Ranges





Predicated Code

- ILP created by IF-conversion leads to an overly conservative IG in Predicated Code
 - might lead to unnecessary interference relationships
 - increase register pressure
- We used Elcor's PQS (Predicate Query System)
 - Used in register allocator to determine which groups of predicated operations are disjoint.
 - I.e. predicated upon complementary predicate registers.
 - avoids unnecessary interference in the interference graphs.

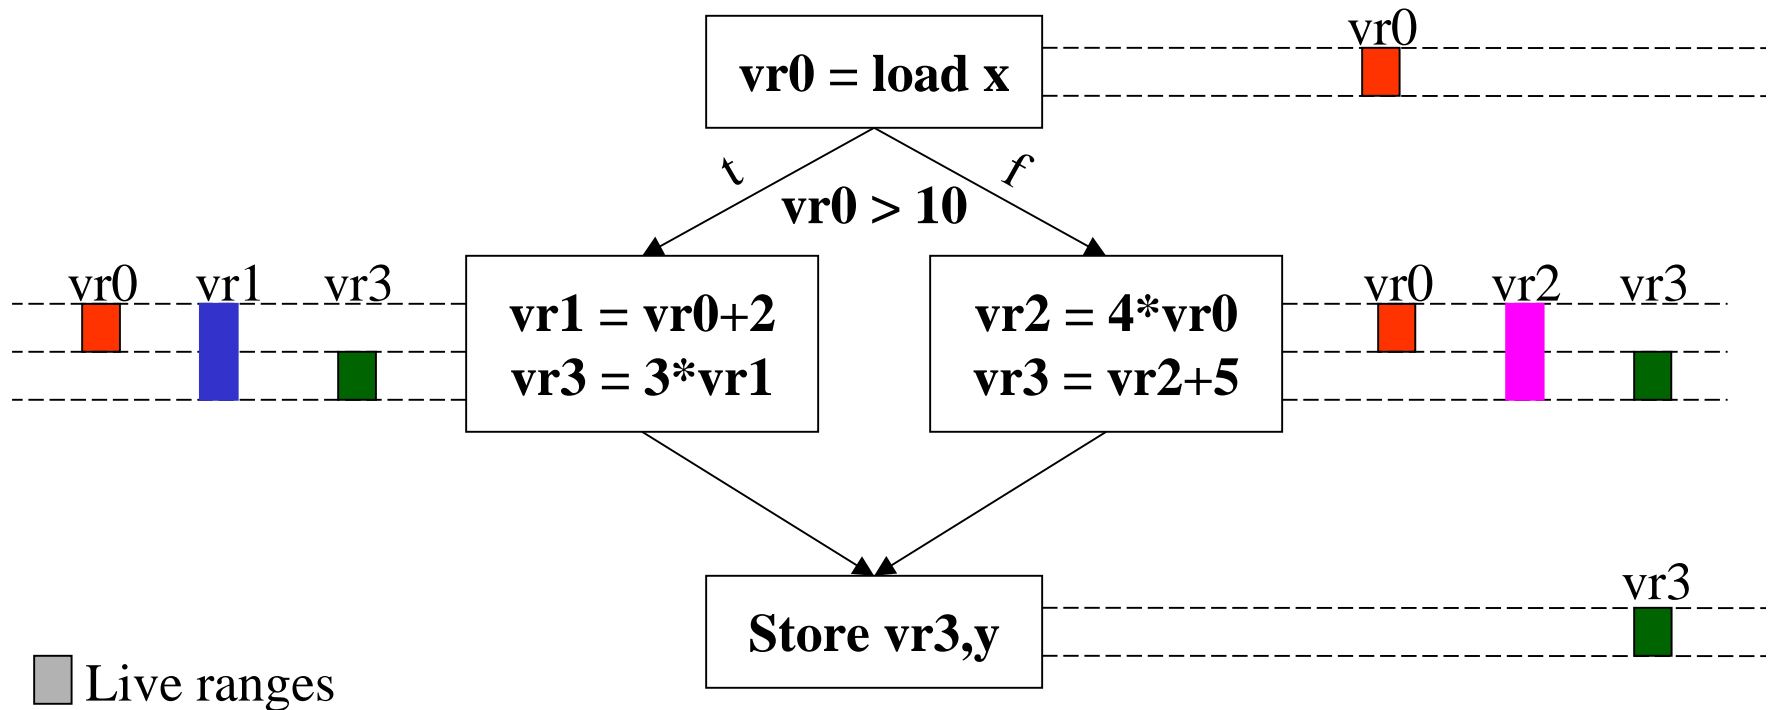


PQS: example

a) Conditional statement


if ($x > 10$) $y = 3 * (x + 2)$; else $y = 4 * x + 5$

b) Flow graph and live ranges for each basic block





Interference Graph






Predicated block and live ranges

bb0-1-2-3

```
vr0 = load x
p1,p2 = (vr0>10)
vr1 =  vr0+2  if p1;  vr2 = 4*vr0  if p2
vr3 =  3*vr1  if p1;  vr3 = vr2+5  if p2
store vr3 ,y
```



 Unconditional
live ranges


 Live range
under p1

 Live range
under p2




Interference Graph

b) Without PQS



c) Using PQS





Register Allocator as a Case Study

- HPL-PD architectural features to be addressed
- Module Interface Issues
 - Input, Output, Module Placement
- Implementation Issues
 - Data structures, libraries, tools
 - How the module was inserted in the compilation path
- Conclusions about the infrastructure



HPL-PD-Specific Allocation Issues

- Predication must be handled correctly
 - support provided by Elcor's Predicate Query System (PQS)
- No other aspects of HPL-PD influenced the design or implementation of the register allocator.
 - VLIW instruction was considered as a sequence of operations.



Module Interface Issues

- Input/Output

- Like all Elcor modules, the register allocator is an IR-to-IR transformer.
- The input is a program graph in the internal IR
 - instructions have been scheduled
 - register operands are virtual.
- Output is program graph in the internal IR
 - register operands are physical registers
 - spill operations have been added to the graph
 - but not yet scheduled



Module Interface (cont)

- Module Placement
 - Register Allocator comes after modulo scheduling and prepass acyclic scheduling and is followed by a postpass acyclic scheduler.
 - Adding a new phase of the compilation process is very easy
 - really just adding a procedure call in the main driver routine.
 - Call is to the main routine of the new module
 - Argument is generally the current procedure being compiled.



Main Elcor Driver Routine


```
Common_process_function(procedure *f)
{
    . . .    //preprocessing, initialization
    . . .    //classic optimizations
              //(copy propagation, etc)
    . . .    //modulo scheduling
    . . .    //prepass acyclic scheduling
    el_solve_reg_alloc(f) // register
                        //allocation
    . . . //postpass scheduling
    . . . //finalization
}
```

We just
added
this call



Implementation Issues

- Data structures used
 - Naturally, the region, operand, and edge classes of the IR were used.
 - In our case, only control edges
 - no register or memory dependence
 - Also a rich class library
 - Hash maps, bit vectors, lists, doubly-linked lists, hash sets, sorted lists, tuples
- Library Routines
 - We used an existing procedure for computing liveness.
 - Also the PQS routines.



Implementation Issues (cont)

- Debugging Support

- Simulator for functional debugging.
- It's easy to examine the Rebel text before and after a phase.

- Modify common_process_function

```
. . .  
ir_write(out,f)  
el_solve_reg_alloc(f)  
ir_write(out,f)
```

```
. . .
```

- Class Browser

- DaVinci tool for examining the Rebel IR in graphical form.



Conclusions

- Development Time
 - 2 person-months implementation time + 1 person-month testing and debugging
 - Once familiar with infrastructure (several more months)
 - Very short development time for a real register allocator in a serious compiler.



Conclusions (cont)

- Value of Infrastructure

- Primary value lay in the framework

- IR, module interfaces, existing modules

- Tools were useful

- PQS
 - Liveness Analysis

but would not have been hard to write from scratch.

- Easy for user community to enrich the existing infrastructure.