

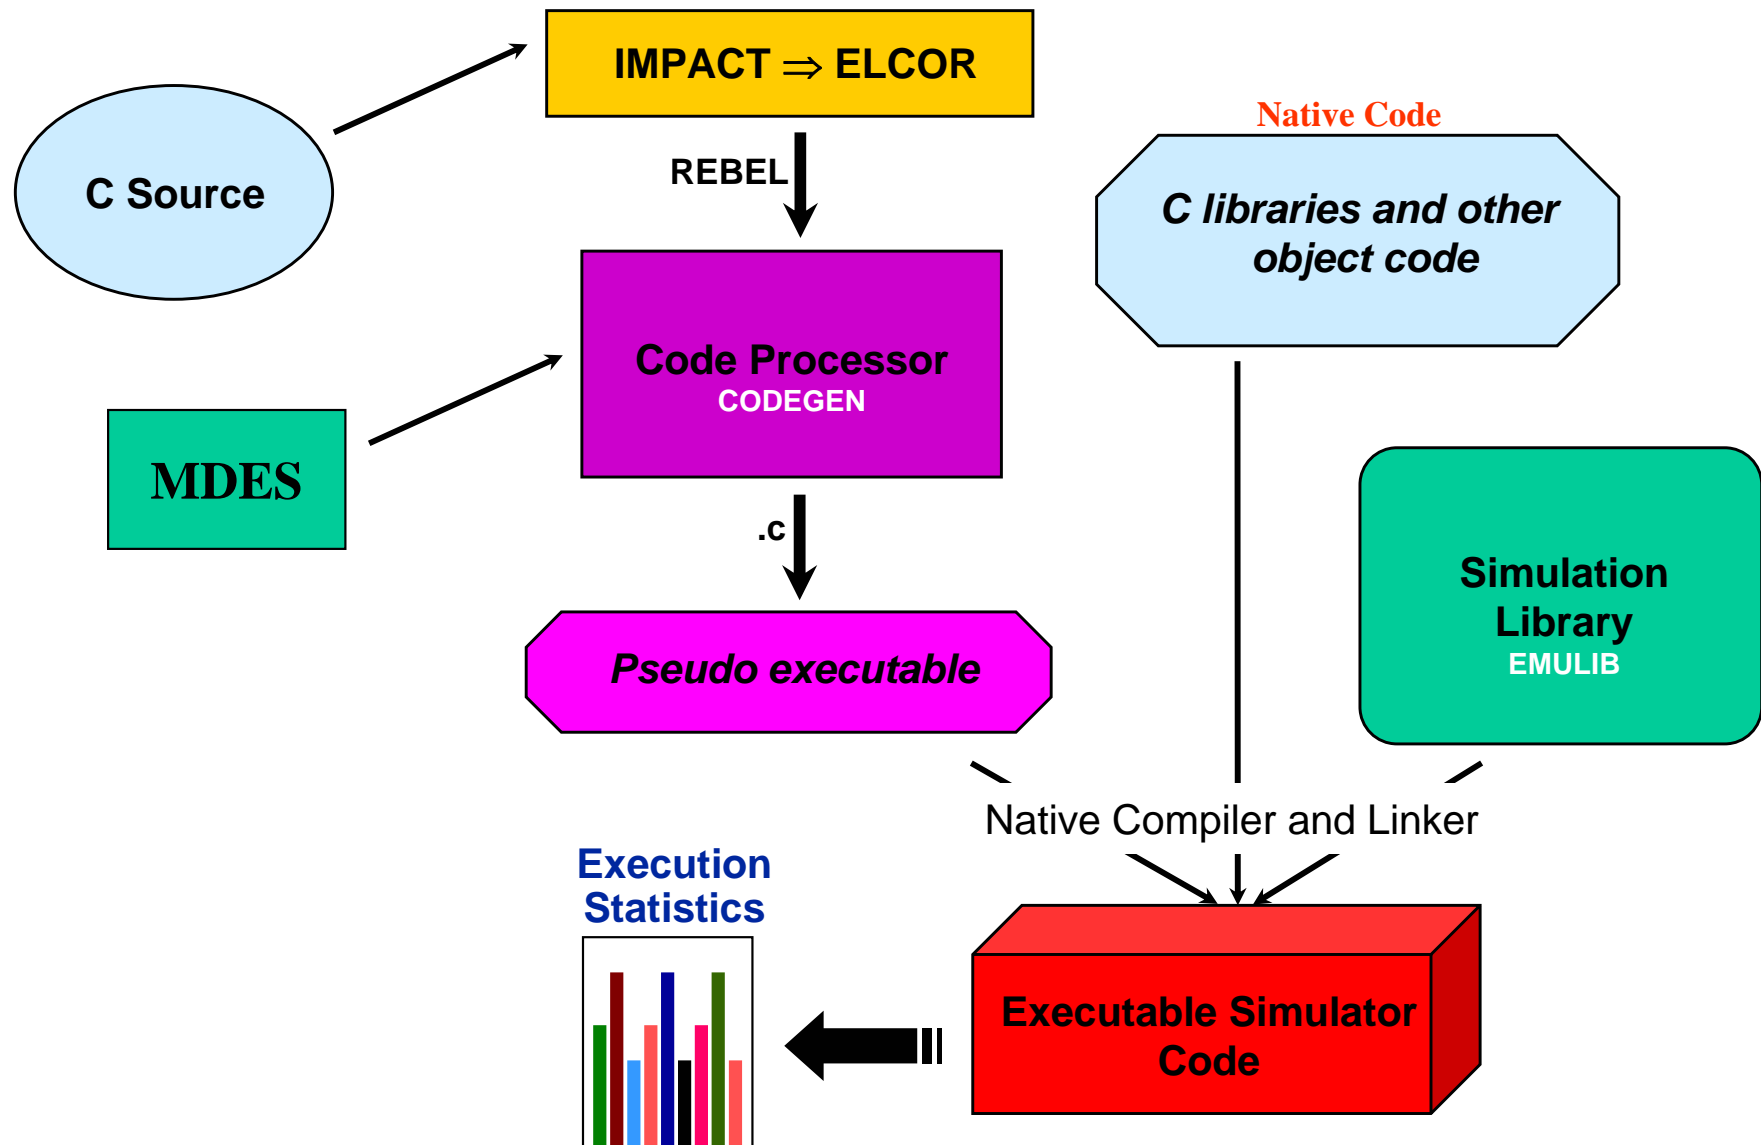


The Trimaran HPL-PD

Simulation Infrastructure

Rodric Rabbah

- The goals of the HPL-PD simulation framework are
 - Emulate the execution of the generated REBEL code on a virtual HPL-PD processor
 - Have the ability to adapt to changes in the machine description
 - Generate *accurate* run-time information
 - Execution clock cycles
 - Dynamic control flow and call trace
 - Address trace
 - Average number of operations executed per cycle





General Description

- The code processor translates the REBEL input to a C equivalent
- The output of CODEGEN is a collection of *pseudo-executable* low-level C files
- A native compiler (i.e. GCC) is used to generate the equivalent machine code

General Description (continued)

- The pseudo-executables are compiled and linked with the *Simulation Library* and any other native libraries or object code
 - The generated executable is the simulator for the specified application
- The resultant executable is run on the host platform to generate statistics and dynamic profile information

EMULIB - The Simulation Back-End



The Simulation Library

- EMULIB Consists of a *virtual machine interpreter* to execute HPL-PD operations
 - Supports EQUALS and LTE execution models
 - Supports speculative execution
- Accumulates various execution statistics
- Generates dynamic information
 - Control flow
 - Operation nullification
 - Address trace

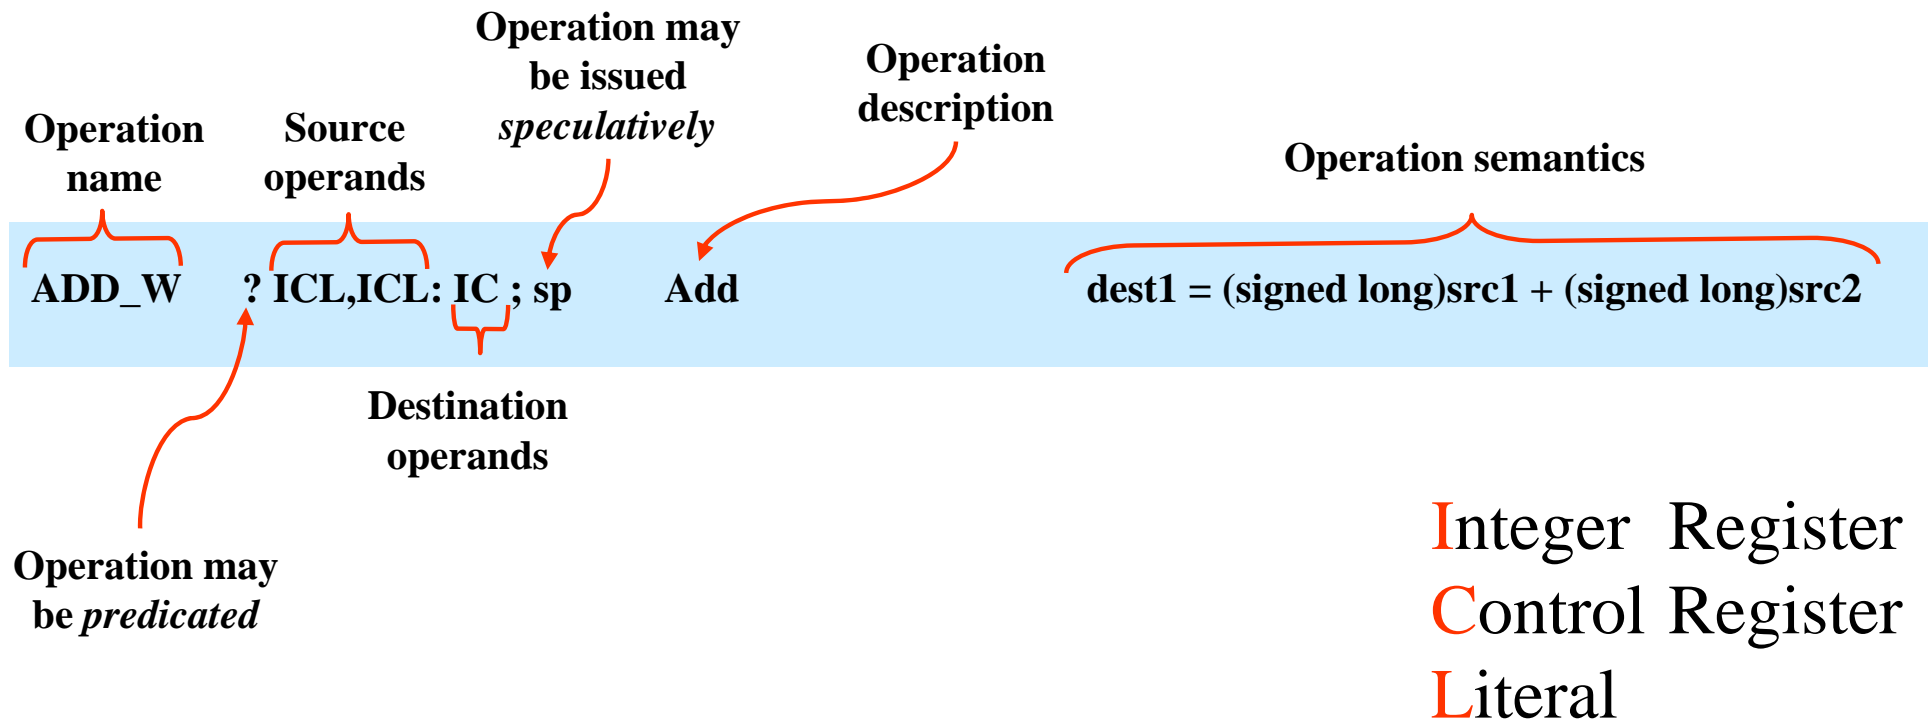


The Simulation Library (continued)

- The virtual machine interpreter (VMI) is invoked on every procedure entry
 - VMI simulates the supplied operation stream in a loop until the procedure returns or the application terminates
- There is one function to simulate every HPL-PD operation
 - The simulation functions are automatically generated from a semantics description file

Simulation Function Generation

- A register transfer level type description of the operation format and its semantics are translated to a C function





Function Generation (continued)

- Special *directives* can be used for special processing
 - `\INIT #define DA_OC(_dest,_pred,_res) if ((_pred) && !(_res)) (_dest) = 1`
instructs the processor to copy the characters following the directive exactly as is
 - `[xxx]`
instructs the processor to redirect the output to a file with name xxx
 - `#` This text is converted to a C style comment
instructs the processor to convert the characters following the directive to a C style comment
 - `\LOOP j X Y Z ...`
instructs the processor to substitute each occurrence of \$j with the X Y Z ...

Function Generation (continued)

- *ops.list* defines all operations of the HPL-PD ISA
- *gen_functions* processes each operation defined in *ops.list* to automatically generate an equivalent C function
 - A C header file declaring the respective prototypes is also generated



HPL-PD Operation Simulation

- The generated functions compute the operation and save the result into local variables
- A computed value is not reflected in the processor state until L cycles after issue time have elapsed
 - L is the operation latency specified by the ELCOR compiler



Example Operation Simulation

```
function void PD_SUB_W_reg_lit(Op *op)
{
    REG r;
    r = (op → src[0] → reg) - (op → src[1] → lit);
    PD_delay(&r, &(op → dest[0] → reg), op → lat[0]);
}
```

- The assignment computes the subtract literal operation and saves the result in local variable *r*
- PD_delay schedules the update of the processor register (op → dest[0] → reg) to occur after (op → lat[0]) cycles

- An operation in the simulation library is represented by a C structure PD_OP in PD.h

```

struct PD_OP {
    unsigned int    Rebel_id;          /* Rebel op id */
    PD_IMPL op;          /* pointer to simulation function */
    PD_PORT         pred;              /* predicate port */

    PD_PORT         src [PD_MAX_SRC];  /* source ports */
    PD_PORT         dest[PD_MAX_DEST]; /* destination ports */
    unsigned int    lat [PD_MAX_DEST]; /* output latencies */

    /* debugging and statistics */
    PD_op_type      type;              /* type of operation */
    unsigned int    Predicated;        /* Set to 1 when Predicated Op */
    unsigned int    Speculated;        /* Set to 1 when Speculative Op */
    unsigned int    Mask;              /* Set to 1 when Mask Op */
};
    
```

- A PD_PORT is an operand
 - Literal and Register operands share the same C structure PD_PORT

```
struct PD_PORT {  
    unsigned int is_reg; /* register or literal operand */  
  
    PD_REG      *file; /* pointer to the register file or the actual literal value */  
    unsigned int num; /* register number when the operand is register */  
    unsigned int rot; /* rotating register number when applicable */  
};
```

HPL-PD Register

- The PD_REG C structure defines the representation of a register in the simulation library
 - One structure is used to represent floating point and integer values
- Each register has a *speculative tag* bit used to support speculative execution

The Simulation Loop

simplified interpreter loop from PD_main.c

```
function void PD_simulate(Op opsTable[])
{
    for (Op *opPtr = opsTable[0]; !done; opPtr++)
    {
        /* perform the operation */
        *(opPtr → op)(opPtr);
    }
}
```

Invokes the function which will simulate the current operation

Pointer to the current operation – contains source and destination operand information, latency information, and guarding predicate operand

The *OpsTable*

- The *OpsTable* is an array of PD_OPs
 - Generated by CODEGEN during the REBEL to C translation

```
PD_OP opsTable_fir[75] = {  
    {  
        70,  
        PD_SUB_W_lit_reg,  
        { {(PD_REG *) 32,0,0x0,0}, {(PD_REG *) PD_GPR_FILE,4,0x0,1}, {0x0,0,0x0,0}, {0x0,0,0x0,0} },  
        { {(PD_REG *) PD_LC,0,0x0,1}, {0x0,0,0x0,0}, {0x0,0,0x0,0}, {0x0,0,0x0,0} },  
        {0,0,0,0},  
        {0,0,0,0},  
        PD_ialu,  
        0,0,0  
    },  
    {  
        ...  
    }  
}
```

The Simulation Loop (continued)

- All operations with the same schedule time are issued without updating the simulator cycle clock
- A MultiOp is a bundle of operations to be issued within the same cycle
 - VLIW or EPIC bundle
- A MultiOp separator is inserted by CODEGEN
 - PD_ACLOCK
 - Commits all actions pending in the action queue
 - Advances the simulation cycle count

The Simulation Loop (continued)

- Every simulated operation uses the PD_delay function to *queue* the pending *action* in an *action queue*
- An action is either a *write* action or a *call* action
 - Write actions copy a value from the simulator's local state to the processor state and is implemented as an assignment
 - Call actions transfers control to the specified function or block
- Note that actions that read values from the processor state for input operands are performed at issue time
 - i.e. Register and Memory reads

Branch Operations

- A call action may be the result of a *branch* or a *branch and link* operation
 - The source register is a pointer to a *block* or a *vector* associated with the target
- When branching to a block, the program counter is updated with the target block address
- When branching to a function, the vector transfers control to a new function
 - Function is either an *external* function or a *wrapper* function
 - external functions are those not processed by CODEGEN
 - Standard library function calls (i.e. printf)
 - wrapper functions are those processed by CODEGEN
 - Wrapper function will eventually transfer control back to the VMI

Speculative Execution Support

- To support speculative execution, a function for intercepting *exceptions* is defined
 - When an exception is raised, the (REBEL) id of the offending operation is recorded but the exception suppressed
- Exception is raised when the speculative bit on any of the source operands is set to true *and* the current operation is not issued speculatively
 - CODEGEN flags speculative operations



Speculative Execution Support (continued)

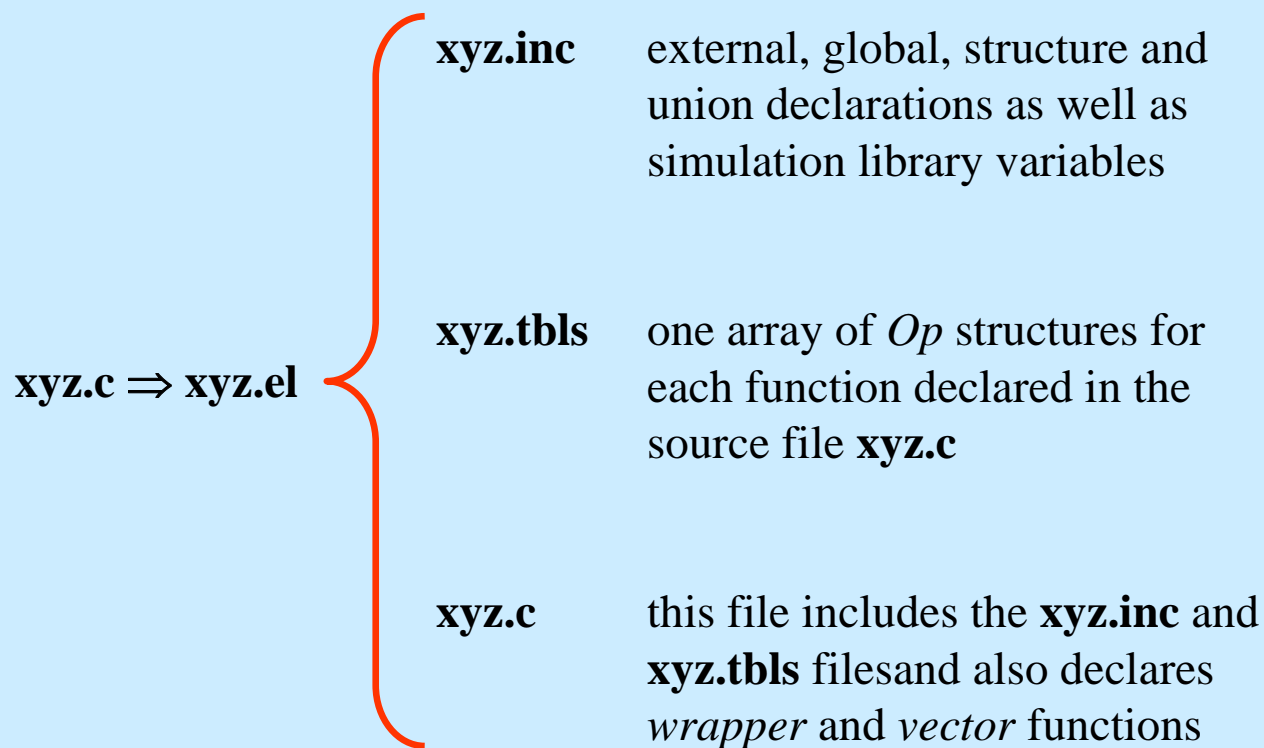
```
function void PD_simulate(Op opsTable[])
{
    for (Op *opPtr = opsTable[0]; !done; opPtr++)
    {
        if (!(opPtr->Mask | opPtr->Speculated)) {
            /* if one of the sources was the result of a speculated instruction
             * that caused an exception, raise the exception now
             */
            if (PD_has_speculated_sources(opPtr)) {
                PD_raise_exception(opPtr->Rebel_id, 0);
            }
            else {
                /* store the current operation's id in case of a delayed exception */
                PD_DO_NOT_TRAP_RAISED_EXCEPTION = opPtr->Rebel_id + 1;
            }
        }
        else {
            PD_DO_NOT_TRAP_RAISED_EXCEPTION = 0;
        }

        /* perform the operation */
        *(opPtr → op)(opPtr);
    }
}
```

CODEGEN – The Simulation Front-End

CODEGEN Overview

- For each application source file specified a set of low-level C files are generated



A `benchmark_data_init.simu.c` file is also generated for the entire application. It contains the definition of a function that initializes the global data for this application.

CODEGEN Overview

- A REBEL input file consists of multiple DATA and CODE sections
- Each DATA section is translated to corresponding C global data declarations
- Each operation in the CODE section is translated to an equivalent pseudo-operation
 - Branch operation require special translation
 - Instrumentation operations and MultiOp separators are inserted during the translation



CODEGEN Components

- DATA processor
 - Declares global variable and constants extracted from source file
- CODE processor
 - Translates a REBEL operation to an equivalent PD_OP
 - Inserts MultiOp separators and instrumentation code
- Vectorizer
 - Generates vector functions
- Packager
 - Generates wrapper functions
- Integrator/Emitter
 - Generates the pseudo-executables

- The REBEL data section contains variable names and layout information
 - Extracted from the host compiler preprocessor

```
data (
  (align 8 _$$$__init_fpc)
  (reserve 16)
  (wf (add (1 _$$$__init_fpc)(i 0)) (f 32))
  (wf2 (add (1 _$$$__init_fpc)(i 8)) (f2 2))
  (global _x (global) (float) (arr (i 128)))
  (align 4 _x)
  (element_size 4 _x)
  (reserve 512)
  (global _pi (global) (float))
  (float 1 _pi (f 3.1415901))
)
```

Global Variable `_x`
Array of 128 bytes

Floating Point Constant

Requires 16 bytes

Write Float Single Precision
Offset 0

Initialize to 32

Global Variable `_pi`

Initialized to value 3.1415901

- `simu_el_data_processor.cpp` parses and translates the REBEL data sections to C declarations
 - Declares all global data variables and constants
 - Declares local floating point constants
 - Generates the *externs* list
 - Generates the function prototype list
 - Generates a function to initialize any variables if necessary
 - This function is invoked only once when the simulation begins

- The Code Processor translates the Rebel input to a C equivalent
 - Rebel operations are translated to a PD_OP data structure
- Instrumentation code is inserted
- Detects operation boundaries and inserts the ACLOCK operation

Operation Components

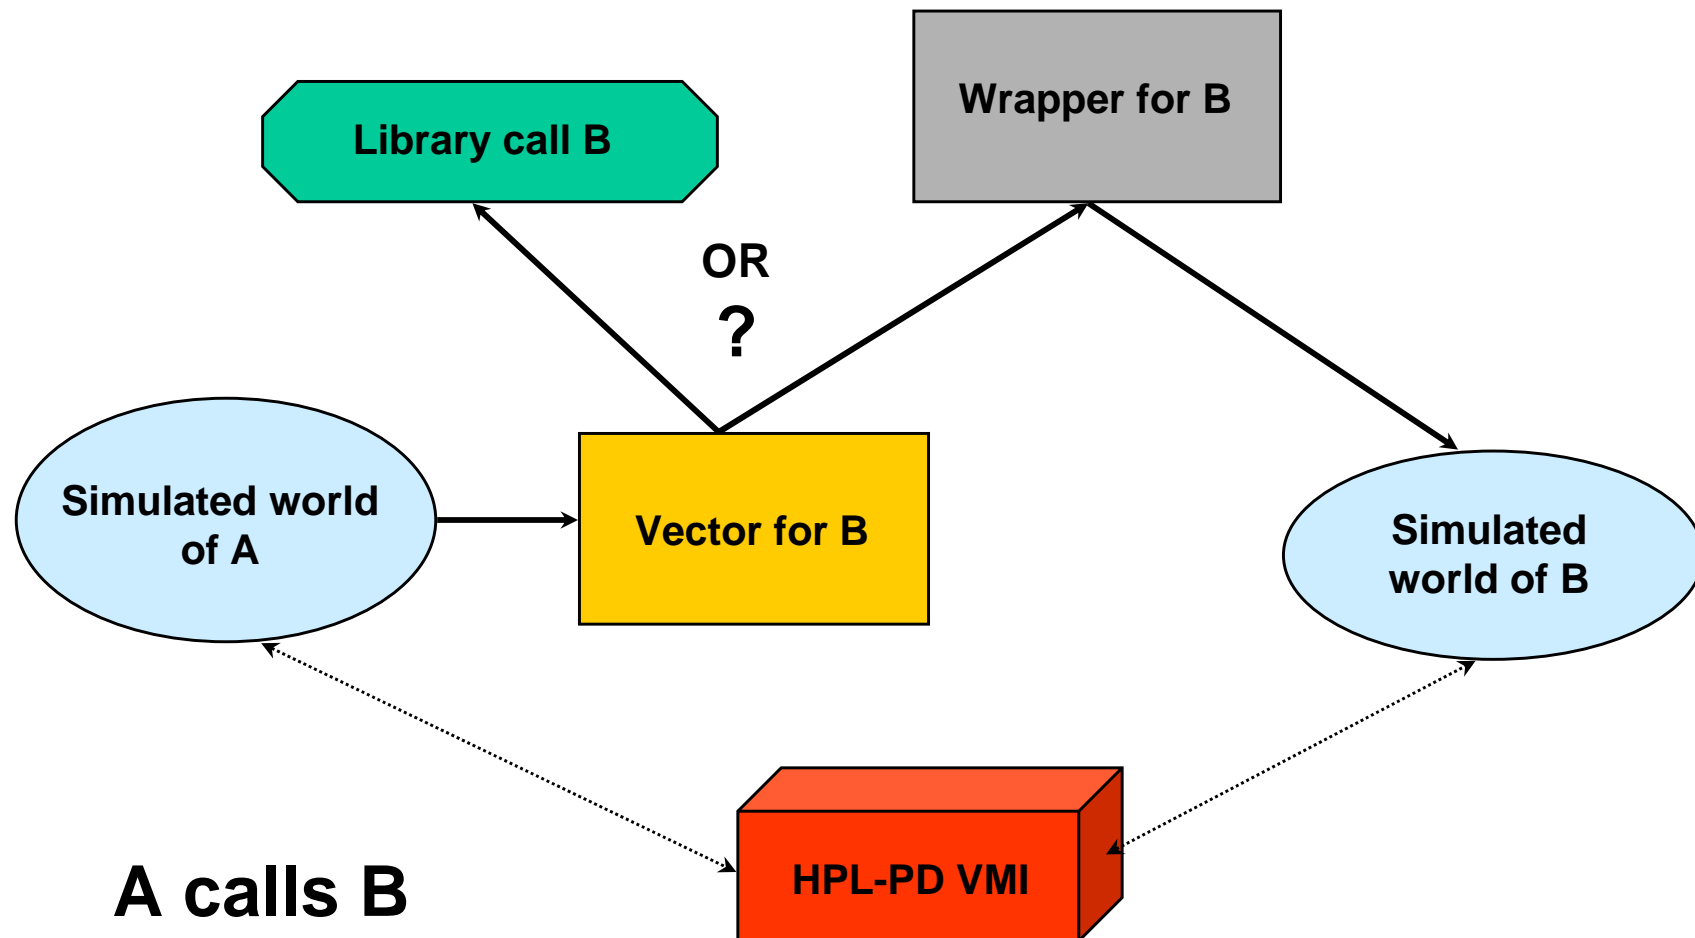
- A REBEL operation consists of
 - Operation id
 - Unique id for information tracking
 - Set of target operands
 - Set of source operands
 - Predicate operand
 - Schedule time
 - OPCODE
 - Information tracking and statistic gathering
 - Attributes (i.e. parameter layout)
 - Flags (i.e. spill operation, speculated operation)

Mapped to equivalent
PD_PORTS

Vector-Wrapper Mechanism

- Allows the simulated code to interact with native C code
 - A *vector* is the only exit interface from the HPL-PD environment
 - A *wrapper* is the only entry interface to the HPL-PD environment
- Calls to external and library functions are made completely transparent
 - Eliminates the need for a linker/loader for the HPL-PD instruction stream
 - Native compiler does the linking

Vector-Wrapper Illustration



- For each function called from the REBEL input file a vector function is generated
 - Transformation code to convert from HPL-PD registers to native function parameters
 - Executes actual call to target function
 - Transformation code to convert the returned function values to HPL-PD registers



Vectorizer Example

C function to which a call is being made

```
int foo ( )
{
    .
    test (p1, p2);
    .
    .
}

int test (int b, double c)
{
    .....
}
```

Vectorized function

```
int _vector_test (void)
{
    *PD_Int_Result =
        test (*PD_Int_Arg_1,
              *PD_Dbl_Arg_1);
}
```

- For every function defined in the input REBEL file a corresponding wrapper is
 - Maps the incoming parameters to the HPL-PD registers and stack
 - Initializes the HPL-PD VMI if necessary
 - Invokes the HPL-PD VMI
 - Upon function return, the HPL-PD registers are mapped to the native parameters

Wrapper Example

The Original C function in source

```
int f( )
{
    .
    test(p1, p2, p3);
    .
    .
}

int test(int b, double c)
{
    ....
}
```

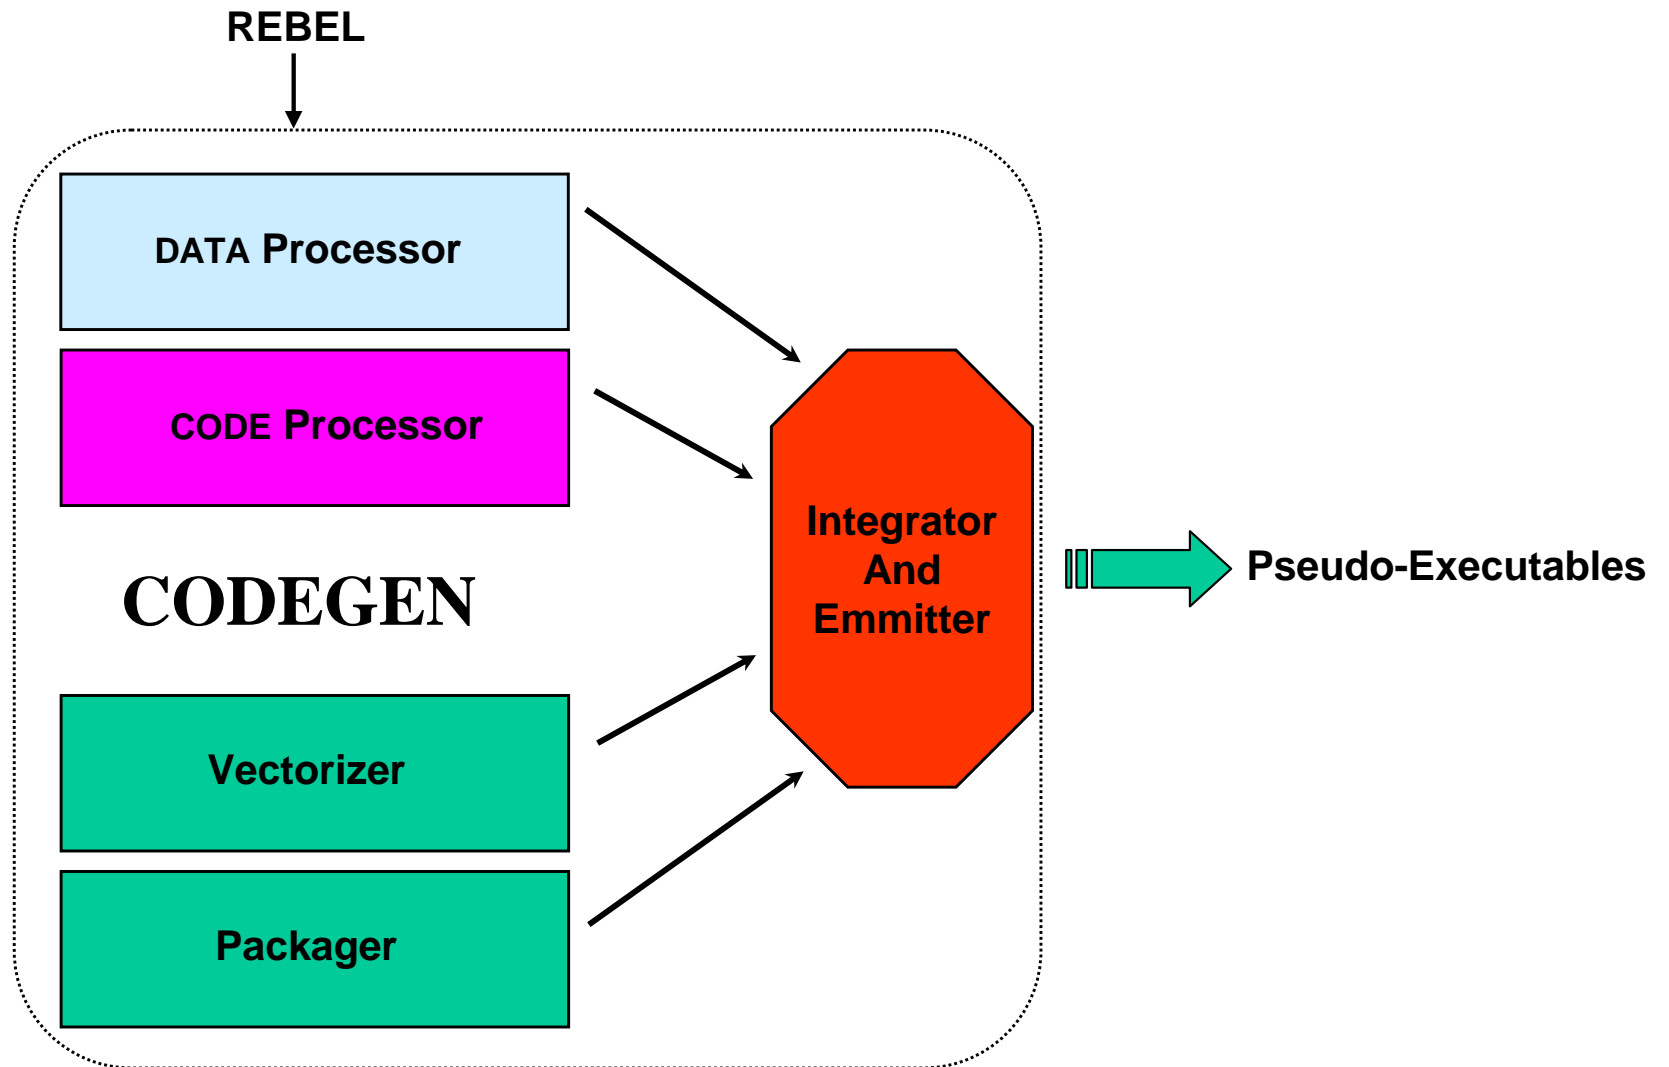
Generated wrapper

```
int test(int b, double c)
{
    *PDS_Int_Arg_1 = b;
    *PDS_Dbl_Arg_1 = c;

    PD_simulate(test_opsTable)

    return *PD_Int_Result;
}
```

Putting It All Together



Integrator/Emitter

- Generates the xyz.c file
 - *#includes* xyz.inc and xyz.tbls
 - Contains the wrapper and vector functions
- Generates the xyz.inc file
 - Includes externs, global variables, floating point constants and global initialization function
- Generates the xyz.tbls file
 - Includes the OpsTables for the functions in xyz.c

Trimaran Extensions and Support



Reconfigurable Cache Simulator

- Smart Memory And Cache Hierarchy Simulator
 - SMACHS
- Memory hierarchy is specified using MDES
- SMACHS can be easily reconfigured and/or extended
 - User may *inherit* and customize cache objects
 - Default behaviors are easy to override



SMACHS Functionality Overview

- Accurate fine-grain cycle-by-cycle simulation
- Simulates DATA or INSTRUCTION cache
 - instruction cache simulation requires instruction format specification
- Parametric
 - uses HMDES to describe structure and behavior
 - user may choose to replace existing data management policies



Functionality Overview (continued)

- Fine-grain specification of a memory unit
 - Type of memory
 - Main memory or cache memory
 - Specify connectivity
 - Memory hierarchy layout, port and bus width
 - Direct mapped or N-way associative
 - Size specification
 - Capacity and line size
 - Transaction latencies
 - Includes port and bus latencies
 - Behavioral aspects
 - e.g. write-thru or write-back



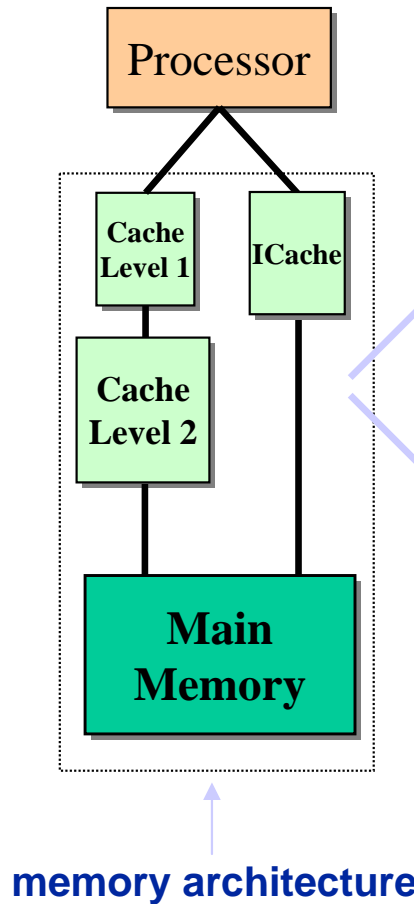
Functionality Overview (continued)

- Simulates complex cache interactions
 - Multiple cache levels, including multiple caches on the same level
- Multiple modes of operation
 - Trace-driven simulation
 - Direct interface to HPL-PD simulator
- Documentation
 - Tech-Report in progress

SMACHS - GUI Support

- Integrated into the Trimaran Laboratory
- Functionality includes
 - Memory hierarchy specification
 - Creating, editing and compiling
 - Trace selection
 - Off-line simulation
 - Statistical analysis
 - General statistics include total computation cycles, stall cycles, number of cache hits and misses
 - Operation specific statistics report the hit and miss ratios as well as the average latencies in delivery data for specific operations

Memory Hierarchy and HMDES



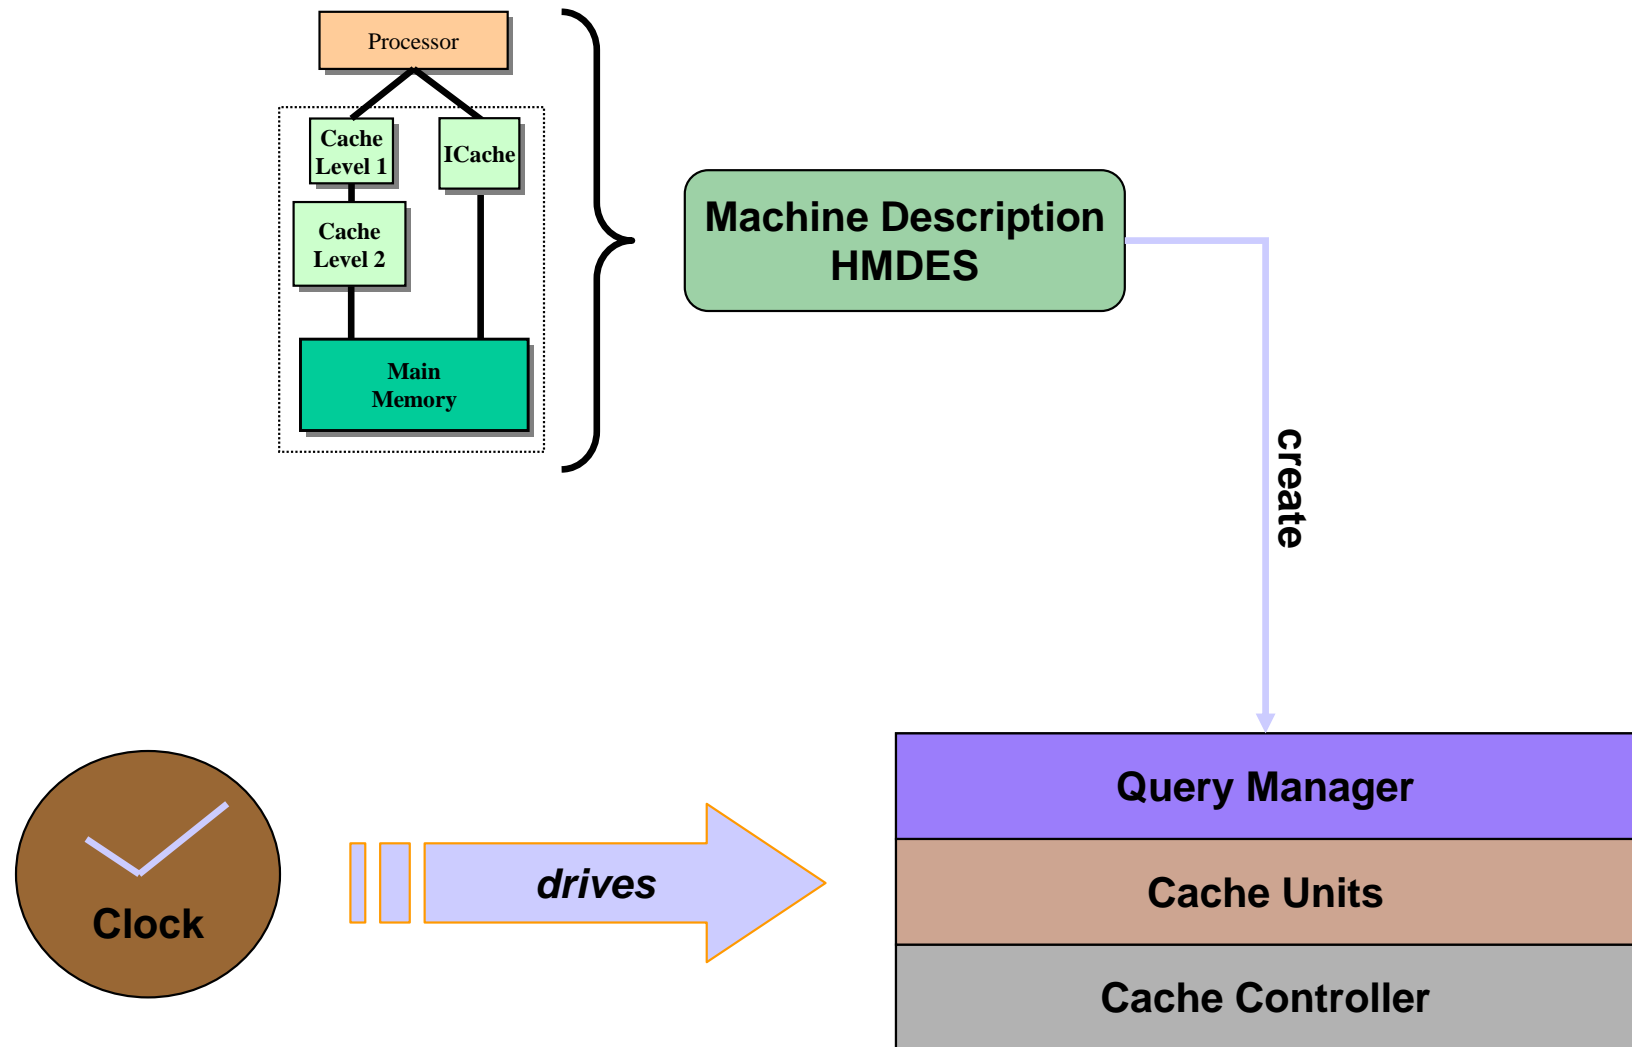
SECTION Cache

```
{
  ...
  // L1Dcache
  L1D(type(Basic_Cache)
    num_lines(64)
    line_size(32)
    associativity(${assoc_direct_mapped})
    replacement(LRU)
    load_hit_latency(1)
    load_miss_latency(1)
    store_hit_latency(1)
    store_miss_latency(1)
    turnaround_latency(0)
    max_misses(1));
  ...
}
```

SECTION Port


```
{
  L2toMainPort(width(32) latency(5) up(L2) down(MAIN));
  L1DtoL2Port(width(16) latency(2) up(L1D) down(L2));
  L1ItoMainPort(width(1) latency(1) up(L1I) down(L2));
  MMUtoL1DPort(width(1) latency(0) up(MMU) down(L1D));
  MMUtoL1IPort(width(1) latency(1) up(MMU) down(L1I));
}
```

SMACHS Design Overview




Query Manager - QM

- Interface for all memory requests
- Simulator or Trace driven
- Multiple output ports



- parametric - width, latency

- Parametric cache units
 - Line size, number of lines, associativity, latencies, replacement policy, etc...
 - Cache-Hash attribute
 - Multiple input/output ports



- Specifies complex behavior
 - Next-Hop-Table as an example
 - Central request routing table
 - Can be used to dynamically dispatch a request along different routes

